

New Systems and Algorithms for Preserving Big-data Privacy in Clouds

Abstract:

In this big-data era, many application vendors (e.g., Uber) store data on clouds. Meanwhile, application vendors and third-parties implement self-defined queries (e.g., MapReduce) to process data, causing two major privacy problems. First, third-parties can leak sensitive fields in data records (e.g., credit cards in Uber orders) through query results. Existing techniques (e.g., differential privacy) add random noise on results, but they often add excessive noise and make results inaccurate. Second, careless/malicious cloud providers can observe the data being queried.

This proposal aims to preserve big-data privacy with a holistic methodology: people still run unmodified big-data queries, and this proposal will automatically prevent sensitive data leakage by accomplishing three objectives.

First, to confine malicious third-parties, we will build Kakute, the first Data Flow Tracking (DFT) system for big-data. Kakute provides easy-to-use APIs for application vendors to tag sensitive data fields, it then automatically tracks unmodified queries and prevents sensitive data flowing to query results. A challenge in existing DFT systems is that propagating tags in data-intensive computations is too slow (e.g., a notable DFT system incurs 128X performance overhead compared to native, insecure queries). By leveraging subtle efficiency features of big-data queries, we will create two fast tag propagation techniques. Our Kakute preliminary prototype presented in [ACSAC '17] incurs merely 32.3% performance overhead.

Second, we will create a Fine-grained Differential Privacy (FDP) technique and its new algorithms. Kakute and differential privacy are synergistic on confining malicious third-parties, because differential privacy allows aggregation computations on sensitive data while hiding individual information. Unfortunately, existing differential privacy techniques are coarse-grained (inaccurate): they often conservatively add excessive noise to all query results, because they can not track which sensitive inputs flowed to which results. By leveraging Kakute, our FDP technique automatically adds noise to only the sensitive (tagged) parts of results, effectively hiding sensitive data and making most results accurate.

Third, to confine malicious cloud providers, we will leverage the Intel SGX security hardware to build the first privacy-preserving compiler for unmodified big-data queries. Existing SGX-based systems for big-data have two challenges: they have to rewrite the queries from Java to SGX-compatible C++, or their trusted components running in SGX are too large (e.g., an entire JVM). Our compiler tackles these challenges by running only the Java big-data queries in SGX using a thin, verified just-in-time translator; it also carries our new, fast SGX-based techniques.

The success of this proposal will effectively preserve big-data privacy in clouds and benefit all people.

Long term impact:

The big-data and cloud computing trends enable great opportunities to all entities, including data providers (e.g., application vendors and computer users), cloud providers (e.g., Amazon), and computation providers (e.g., application vendors and their third-party partners).

Unfortunately, despite decades of effort, data leakage remains one of the most severe threats in clouds. In the perspective of data providers (owners), severe data leakage incidents have been triggered by both computation providers (e.g., some iCloud third-parties leaked celebrity accounts on Internet in 2017) and cloud providers (e.g., the 2013 Yahoo Cloud compromise and the 2014 J.P. Morgan account leakage).

In the perspective of big-data queries, data leakage have happened both from within and outside. When a query runs, the within-code can often be buggy or malicious. Cloud providers, which locate outside the queries, have also incurred numerous data leakage incidents due to compromises on external attacks and insider attacks. Therefore, this proposal takes a holistic methodology: preventing leakage for data providers from both within (Objective 1 and 2) and outside (Objective 3) the queries.

In the short term, to confine malicious third-parties in private clouds (i.e., a data provider is the cloud provider), we plan to accomplish both Objective 1 and Objective 2. Objective 1 proposes Kakute, the first Data Flow Tracking (DFT) system for big data queries. Kakute tackles a notorious performance challenge on porting DFT to data-intensive computations. To achieve a robust DFT architecture for distributed big-data frameworks (e.g., Spark), Kakute completely captures the frameworks' inter-computer data flows. We have implemented a Kakute prototype and integrated it with Spark. Kakute carries built-in checkers for four security and reliability problems: sensitive data leakage, data provenance, programming bugs, and performance bugs. Kakute incurs a moderate performance overhead of 32.3% compared to native, insecure queries. It also effectively detects 13 real-world security and performance bugs. These promising preliminary results have been presented in [ACSAC '17] and [TPDS '17].

Kakute and differential privacy are synergistic on confining malicious third-parties: Kakute enforces mandatory access control on sensitive data, but it may cause some query results containing sensitive data to be missing; differential privacy allows the aggregation computations on sensitive data while hiding individual privacy, but due to the lack of precisely tracking data flow, it often suffers from excessive noise and inaccurate results. Therefore, the Objective 2 of this proposal takes the first significant step to integrate DFT and differential privacy and creates a Fine-grained Differential Privacy (FDP) technique: by leveraging Kakute, FDP automatically adds noise to only the sensitive parts of results, effectively hiding sensitive data and making most results accurate.

In the intermediate term, we plan to confine malicious public cloud providers by accomplishing Objective 3. The Intel SGX hardware enforces strong confidentiality for data and code even if the cloud is malicious. Moreover, SGX is a good fit for big-data queries because these queries do data-intensive computations in userspace and rarely invoke system calls.

However, existing SGX-based systems for big-data have two major challenges: they have to manually rewrite the Java big-data queries into SGX-compatible C++, or their trusted computing base (TCB) is too large (e.g., an entire JVM). To tackle these two challenges, Objective 3 will build the first privacy-preserving Java compiler with minimum TCB. Our new compiler will run only the Java big-data queries in SGX with a thin, verified just-in-time translator, and the rest of JVM is outside SGX without affecting the privacy of the data being queried.

In the long term, by integrating the outcomes of all the three objectives in this proposal and extensively applying them on real-world software, we will help data providers enforce comprehensive privacy against both computation providers and cloud providers. This will benefit almost all people. For instance, many Hong Kong finance companies demand strong privacy for their data deployed in clouds. Moreover, we envision that the outcomes of this proposal will broadly promote other security techniques (e.g., software integrity and availability).

Objectives:**1. [To build the first Data Flow Tracking (DFT) system for private clouds]**

We will create Kakute, a fast DFT system that can track and prevent sensitive data leakage in self-defined big-data queries. We will make Kakute support diverse big-data queries on large, popular datasets, and we will make Kakute incur reasonable performance overhead compared to native, insecure queries.

2. [To create a Fine-grained Differential Privacy (FDP) technique for private clouds]

We will leverage Kakute to develop FDP and its new algorithm, which will only add noise to sensitive data fields or query results, preserving strong differential privacy for sensitive data and good accuracy for most results. We will extensively study FDP's accuracy improvements on both sensitive and insensitive data compared to existing differential privacy techniques.

3. [To construct the first compiler for big-data privacy in public clouds]

Our compiler will support unmodified big-data frameworks by creating a thin, verified translator that automatically translates Java bytecode to SGX-compatible code. We will quantitatively evaluate the performance overhead of our compiler and whether it can protect data privacy against real-world privileged attacks.

1 Research Background

This proposal has three entities: data providers, cloud providers, and computation providers (who write big-data queries). In private clouds, data providers are cloud providers; in public clouds, they differ. This section shows relevant techniques (§1.1, §1.2, and §1.3), motivation (§1.4), and related work (§1.5 and §1.6).

1.1 Big-data computing frameworks

Big-data frameworks (e.g., Spark [74] and MapReduce [16]) are popular for computations on tremendous amounts of data records. These frameworks provide self-defined Java functions (e.g., `map/reduce`) to let computation providers write their algorithms, and the frameworks automatically apply these functions on the data stored across computers in parallel.

To avoid excessive computation, big-data frameworks adopt a *lazy transformation* approach [47, 73, 74]. Spark often uses lazy transformations (e.g., `map`), and calls to these transformations only create a data structure called RDD with *lineage* (the sequence of transformations on data records). Actual transformations are only triggered when materialization operations (e.g., `collect/count`) are called. Collecting operations trigger transformations only along lineages, so unnecessary computations are avoided. **Objective 1** will leverage lazy transformation to create a fast data flow tracking technique: Reference Propagation (§2.1).

1.2 Software-based privacy techniques

Data Flow Tracking (DFT) is a powerful mandatory access control technique for preventing sensitive information leakage [45]. DFT attaches a tag to a variable (or object), and this tag will propagate during computations on the variable at runtime. DFT is used in various areas, including preventing sensitive data (e.g. contacts) leakage in cellphones (TaintDroid [21]), web services [50], and server programs [36]. No DFT system exists for big-data computing, so **Objective 1** (§2.1) will create the first DFT system for big-data.

Complimentary to DFT, statistical techniques, including K-anonymization methods [39, 63] and differential privacy [41, 43, 53], allow the aggregation of sensitive data while adding random noise on inputs or query results to preserve individual privacy. However, statistical techniques are either not secure (K-anonymization) or suffering from great losses of accuracy (differential privacy). Recent work [28] reports more than 30% loss of accuracy. For a query results, low accuracy means bad utility. For instance, a KMeans program will return centroids far from the actual ones, because the accuracy loss rate is much larger than the training error rate (several percents in practice).

A key reason for this bad utility problem is that differential privacy can not track how sensitive data fields flow to query results, so they have to take a coarse-grained approach, which conservatively adds noise to all data fields or all query results. **Objective 2** (§2.2) will propose a novel fine-grained differential privacy technique, which combines the strengths of DFT and differential privacy.

1.3 Hardware-based privacy techniques

Trusted Execution Environment (TEE) is a promising technique for protecting computation in a public cloud even if the cloud’s operating systems and hypervisors compromise. For instance, Intel-SGX [29], a popular commercial TEE product, runs a program in a hardware-protected *enclave*, so code and data are protected from outside. Compared with the approach of computing on encrypted data (§1.5), TEE is 100X to 1000X faster. For instance, a SGX-based system Opaque [76] incurs a moderate performance overhead of 30% compared to native, insecure big-data queries.

However, to practically run Java big-data queries with SGX, two open challenges remain. First, existing SGX-based systems [76] require computation providers to manually rewrite the readily pervasive Java queries into SGX-compatible C++, a time-consuming and error-prone process. Second, existing SGX-based systems for big-data have too large Trusted Computing Base (TCB). Existing systems (e.g., SGX-BigMatrix [55]) run a whole language interpreter (e.g., JVM and Python runtime) in enclaves, causing a too large (and too dangerous) TCB: JVM code comes from many different parties/vendors and extremely hard to verified. **Objective 3** (§2.3) tackles these two open challenges by building a new just-in-time compiler.

1.4 Motivation of objectives

Data leakage (or breach), defined as the leakage of sensitive customer or organization data to unauthorized users [35], is a top security threat [4, 32] in cloud computing. In a data provider’s perspective, both computation providers (e.g., the 2017 iCloud account leakage caused by third-parties [71]) and cloud providers (e.g., the 2013 Yahoo Cloud compromise [64]) have caused severe data leakage and huge financial loss. This proposal aims to preserve the data provider’s privacy by going two directions. First, we will propose two novel complimentary techniques in **Objective 1** (KAKUTE) and **Objective 2** (fine-grained differential privacy) to protect privacy against the computation providers in private clouds. Second, we will propose **Objective 3** (a new privacy-preserving compiler) to protect privacy against the (public) cloud providers. By integrating the outcomes from all three objectives, data privacy will be effectively preserved.

1.5 Related work by others

Computing on encrypted data. Homomorphic encryption [25] is a technique for computing on encrypted data in untrusted environments. Homomorphic encryption contain two kinds: Fully homomorphic encryption (FHE) and partial homomorphic encryption. Partial homomorphic encryption (e.g. Additive Homomorphic Encryption [48]) incurs a much lower overhead compared with FHE. An evaluation [24] on FHE shows a 10^9 X slowdown, unacceptable in practice. Systems that adopts PHE (e.g., Monomi [67], Crypsis [62], CryptDB [52], MrCrypt [65]) reports a much lower overhead, but it has limited expressiveness (e.g., SQL operators) or requires extra trusted servers. Seabed [49] proposes asymmetric encryption schemes and reduces the performance overhead of AHE, but its expressiveness is still quite limited.

SGX-based systems. Intel SGX is a promising technique to provide privacy-preserving analytic in public clouds. Compared with software-based solutions, hardware-based solutions incurs much lower overhead. TrustedDB [7] is a hardware-based secure database. VC3 [54] proposes a secure distributed analytic platform with SGX protection on MapReduce [16]. Opaque [76] supports secure and oblivious SQL operators on SparkSQL [6]. However, all these systems have limited expressiveness (e.g., SQL operators), and VC3 even needs to rewrite the program with C++. A recent work [46] proposes an oblivious machine learning framework on trusted processors. SGX-BigMatrix [55] proposes an oblivious and secure vectorization abstraction on Python, but its trusted components are too big (an entire Python runtime).

Big-data privacy systems. Leakage of sensitive data is a top security threat [5, 35] in clouds because various user sensitive data is stored and processed in clouds. Airavat[53], PINQ [41] and GUPT [43] propose to apply differential privacy [20] in MapReduce in order to prevent leakage from malicious queries, but existing differential privacy techniques often produce inaccurate results. Sedic [75] proposes to offload sensitive computations to private clouds. MrLazy [3] proposes a framework of combining data provenance and static DFT analysis for self-defined queries, to provide fine-grained information flow for security. However, static DFT is not precise and may suffer from false positives. In contrast, our proposed KAKUTE system provides fine-grained DFT and precise mandatory access control on sensitive data at runtime.

1.6 Related work by the PI and co-I

The PI is an expert on secure and reliable distributed systems [11–14, 22, 31, 69, 70, 72]. The PI’s works are published in top conferences on systems (OSDI, SOSP, SOCC, TPDS, and ACSAC) and programming languages (PLDI and ASPLOS). Recently, the PI has collaborated with Huawei to launch a technology transfer project based on his dependable distributed system [70]. The co-I is an expert on high-performance computing [2, 9, 34, 42, 77], fault-tolerance [56, 57], and Java compilers [37, 58, 68]. The co-I’s works are published in top systems conferences (Cluster, SC, and ICPADS) and journals (JPDC, TPDS, IEEE Tran. Computers). As preliminary results for this proposal, the PI has presented KAKUTE [31] in ACSAC ’17, and the PI and co-I have collaborated to present CONFLUENCE [22] in TPDS ’17.

Preliminary results. We have implemented a KAKUTE prototype and evaluated it on six popular big-data algorithms, including three text processing algorithms WordCount [61], WordGrep [38] and TwitterHot [61], two graph algorithms TentativeClosure [61] and ConnectComponent [61], and one analysis program MedicalGroup [51]. We evaluated these algorithms with real-world datasets that are comparable with related systems [10, 26, 30]. Our evaluation shows that: (1) KAKUTE incurred merely 32.3% overhead (Figure 3) with INTEGER tag, about two orders of magnitudes faster than a recent DFT system Phorspor [8]; and (2) KAKUTE effectively detected 13 real-world security and performance bugs presented in other papers [15, 26, 53]. These promising preliminary results have been presented in ACSAC ’17 and TPDS ’17.

Future directions. We will extend KAKUTE in three directions. First, we will port KAKUTE onto more big-data frameworks, including PIG [47] and HADOOP [27]. Second, we will extend KAKUTE to detect broader types of real-world security bugs. Third, we will apply KAKUTE to augment other complementary privacy techniques, including anonymization techniques and differential privacy (**Objective 2**).

2.2 Objective 2: developing the Fine-grained Differential Privacy (FDP) technique

KAKUTE (**Objective 1**) strictly prevents sensitive data flowing to IO functions or query results, but in some scenarios it is still desirable to let computation providers acquire aggregation results (e.g., the sum of citizens who have got cancer in a country) on sensitive fields as long as individual information is not leaked. Differential privacy [17, 18] can enforce statistical bounds on aggregation results and prevent individual information leakage, so it is complementary to DFT and has attracted much attention recently.

2.2.1 Challenge: existing differential privacy techniques are coarse-grained and thus inaccurate

Existing differential privacy techniques often suffer from low accuracy for query results. To prevent computation providers revealing individual data, differential privacy typically adds noise either on input data records or query results. However, due to the lack of precisely tracking how inputs are computed and propagated to outputs, to enforce statistical guarantee on outputs, differential privacy often conservatively add the same noise to all fields of a data record and to all records, causing inaccurate results. For instance, prior work [28] reports more than 30% loss of accuracy when the security guarantee is high (the probability of leakage is low). Therefore, a KMeans program will return centroids far from the accurate ones. This low accuracy makes results useless as it is much larger than the KMeans training error rate (a few percents).

2.2.2 FDP and its new algorithm

Our key insight is that DFT and differential privacy can complement each other, getting the best of both worlds. Considering each data record, DFT can precisely track how sensitive data fields flow to which query result, so differential privacy needs only add noise to the sensitive input fields or results. Considering all data records, DFT can also distinguish which records have higher security sensitivity levels, then we can add more noise accordingly.

This insight nurtures Fine-grained Differential Privacy (FDP). In FDP, each record belongs to a user who assigns a security tag to all its data records. A security tag that is related to the privacy budget ϵ (or accuracy level). When the privacy budget is high, the probability of leakages is high. We adopt the personalized differential privacy model in recent work [33]. $\Phi(x)$ return the ϵ of a particular record x . The threat model of FDP is the same as KAKUTE’s (Figure 1), because FDP aims to defend against malicious computation providers in private clouds.




Figure 3: KAKUTE execution time normalized to native Spark executions. 100% means no overhead.

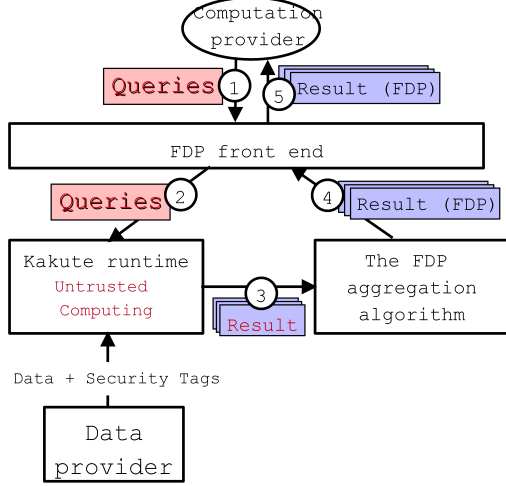


Figure 4: The workflow of FDP with five steps.

Definition 2.1. *Differential Privacy* For two neighbor dataset D and D' differing at record x , a mechanism $\mathcal{M}(y)$ is differentially private with the following condition:

$$\Pr[\mathcal{M}(D) \in O] \leq e^{\Phi(x)} \times \Pr[\mathcal{M}(D') \in O] \quad (1)$$

Intuitively, Differential Privacy guarantees that the probability of producing different result with neighboring dataset is low. For the personalized model, the probability is different for records belonging to different users, so that different users have various levels of protections.

To start with, we need to determine the relation of the privacy budget and the security level (a `getTagLevel` function). We adopt a deterministic model, and there are 6 security levels: insecure, dp_1, dp_2, dp_3, dp_4 and non-released. Insecure records can be release directly, while non-released can not be used in any computation to the final result. dp_1 to dp_4 varies in terms of their privacy budgets for differential privacy.

Theorem 2.1. *Laplace Mechanism [17] Adding noise with Laplace distribution $\text{Laplace}(\frac{\Delta f}{\epsilon})$ enforces Differential Privacy, and global sensitivity Δf is defined as*

$$\Delta f = \max \|f(D) - f(D')\|_1 \quad (2)$$

To enforce differential privacy, one approach is to use the ϵ inferred by the highest security level of each dimension and to add noise to the output directly, but this approach is too naive for the diversity of security levels. Instead, we adopt the sample-and-aggregate approach in previous work [60]. In this approach, data is partitioned into multiple parts (each part has a size of m). Each partition can have different sensitivity levels. The aggregator adds noise to the result of each partition according to its security level. Data is partitioned into multiple parts denoted as p_1, \dots, p_n . The security level and its corresponding ϵ are $\epsilon_1, \dots, \epsilon_n$.

Theorem 2.2. *For any output record in O , the estimator is $\max_{i \leq k}(\epsilon_i)$ -differentially private.*

Proof. For each dimension, we can divide the output dataset D as k disjoint parts D_1, D_2, \dots, D_k according to their security levels. According to [59], each $f(D_i) (i \leq k)$ is ϵ_k -differentially private. For each D_i , we have $\Pr[f(D_i) \in S] \leq e^{\epsilon_i} \Pr[f(D'_i) \in S]$, suppose $\epsilon_{\max} = \max_{i \leq k}(\epsilon_i)$,

$$\begin{aligned} \Pr[f(D)] &= \Pr[f(D_1 + D_2 + \dots + D_k)] \\ &= e^{\epsilon_1} \Pr[f(D'_1)] + \dots + e^{\epsilon_k} \Pr[f(D'_k)] \\ &\leq e^{\epsilon_{\max}} (\Pr[f(D'_1)] + \dots + \Pr[f(D'_k)]) \\ &= e^{\epsilon_{\max}} \Pr[f(D')] \end{aligned} \quad (3)$$

Therefore, $f(D)$ is $\max_{i \leq k}(\epsilon_i)$ -differentially private. \square

Algorithm 1: The FDP aggregation algorithm

Input: Dataset T , dataset size N , privacy budget ϵ_k
 for security level k , output range (min, max)
 $n =$ a suitable partition
for $i \leftarrow 1$ **to** n **do**
 $O_i \leftarrow f(T_i)$;
 if $O_i > \max$, $O_i \leftarrow \max$ if $O_i < \min$, $O_i \leftarrow \min$
for dimension j of the output O **do**
 $k \leftarrow \text{getTagLevel}(O_j)$;
 $O_j \leftarrow \frac{1}{n} \sum_{i=1}^n O_{ij} + \text{Lap}(\frac{\max - \min}{n\epsilon_k})$
Output: O

In the above aggregation algorithm, the error of the final result comes from two parts: the Laplace noise error and the partition error. It is crucial to reduce the final error incurred by this algorithm while keeping the differential security guarantee. We can adopt a hill-climbing approach (future work).

Future directions. We will fully develop this FDP technique by going along three directions. First, we will continue to optimize the algorithm and reduce its final error rate. Second, we will do an extensive study on real-world big-data queries and quantify the improvements on result accuracy. Third, currently we propose a end-to-end differentially private computation system, which adopts a personalized differential privacy model [33] and improves usability without reducing the security guarantee. In future explorations, we can consider an even more fine-grained (ϵ, δ) -differentially [19] private model.

2.3 Objective 3: creating a privacy-preserving compiler for big-data queries in public clouds

Recent real-world privacy breaches have shown that sensitive data are often leaked while being processed in public clouds, including clouds compromises on external attacks [71] and insider attacks [4]. Trusted Execution Environment (TEE) is a promising technique to protect computation on public clouds even if the cloud’s operating system is compromised. For example, Intel-SGX [29] runs programs in a enclave, so code and data are protected and can not be seen by the attackers. Meanwhile, SGX is good fit for big-data queries because these queries are data-intensive in userspace and they hardly invoke system calls (OS kernel can easily break SGX’s security on memory). A latest big-data analytic system Opaque [76] reports only 30% overhead compared to native, insecure executions.

2.3.1 Challenges: existing SGX-based systems require rewriting queries and have too-large TCB

Despite recent advances (Opaque [76], VC3 [54], Azure/Coco [1], and SGX-BigMatrix [55]) on building SGX-based systems for big-data, two major challenges remain. First, enclaves require completely rewriting the readily pervasive Java big-data queries into C++ [1, 54, 76], an time-consuming and error-prone process. Second, to easy implementation, existing systems typically run the entire language runtime (e.g., JVM or Python runtime [55]) within SGX, causing the Trusted Computing Base (TCB) to be too large and vulnerable (e.g., JVM has millions of lines of code from many companies and is vulnerable to insider attacks [4]).

2.3.2 MAAT: a just-in-time (JIT), privacy-preserving Java compiler for big-data queries




Figure 5: Threat model of MAAT. Data records with blue color are encrypted, and white color are plaintext. Shaded (grey) components may leak data, and MAAT is designed to defend against them.

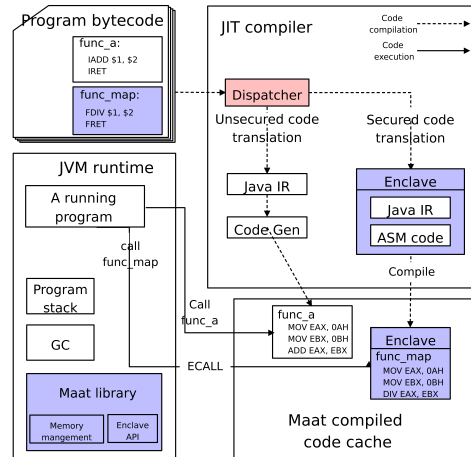


Figure 6: MAAT JIT compiler architecture. Key components are shaded (and in blue).

We propose MAAT, the first compiler that runs unmodified Java big-data queries in SGX enclaves securely with minimal TCB (i.e., the TCB contains only SGX and self-defined code itself). MAAT works as a Java JIT compiler which automatically compiles self-defined big-data functions (e.g., map/reduce) into SGX-compatible assembly instructions. Therefore, the JVM itself does not run in MAAT’s enclaves.

The goal of MAAT is to preserve the confidentiality of data while being processed in public clouds. Other attacks such as changing execution paths (i.e., integrity) have been well defended in prior work [23], and MAAT can directly use it.

Figure 5 shows MAAT’s threat model: both SGX and computation providers are trusted, and cloud providers are malicious. Figure 6 shows the architecture of MAAT. In MAAT, both the translation of self-defined code and the execution of the code are protected by enclaves, so that even if the cloud’s OS compromised, it can not see the executions of big-data queries or inject malicious code into the queries during MAAT’s translation.

Our MAAT architecture contains two software components: our JIT translator and our own management library. The JIT translator is a thin layer which translate a Java bytecode instruction into a number of SGX-compatible assembly instructions. For instance, in Figure 6, an FDIV Java bytecode instruction translates to two MOV and one DIV assembly instructions). The memory management library is for our own use of encryption/decryption on data records and maintaining SGX memory for the queries at runtime. We will proactively implement these two components to be easy to verify (use as few as function recursive calls and loops) as in other verification practice [44], and we plan to use state-of-the-art verification techniques [44] to verify both of them. Then, we do not need to include them in MAAT’s TCB, greatly reducing its TCB.

One subtle performance challenge for MAAT is that it should have reasonable performance overhead compared to native executions. When calling into and out of a function in enclaves, an ECALL and OCALL will be invoked in the CPU and enclave transitions are invoked. Such transitions are several hundreds times slower than user function calls. Moreover, encryption and decryption on data records are invoked during such transitions. Our study on a SGX-based big-data system Opaque [76] shows that it incurs 3.4k enclave transitions for processing only 10k data (with two queries `select` and `groupBy`), which confirms the challenge.

To mitigate this challenge, we will create a new enclave runtime abstraction called Data-locality-aware Asynchronous Enclave calls (DAE). DAE converts the synchronous enclave calls (similar to Java function calls) to asynchronous, data-locality-aware calls into enclaves. Specifically, DAE will run a number of n processes (E_1 to E_n) in an enclave on each computer. When a JVM process P calls a big-data query function, the call and its parameters are appended to a queue to DAE, and DAE arranges a process E_i with good data locality (e.g., according to prior arrangements and the decrypted data held by E_i) to execute the call. The call result is appended to a return queue of the DAE for process P . We expect that DAE will achieve reasonable performance, data locality, and parallelism.

Future directions. By realizing a privacy-preserving Java JIT compiler for public clouds, MAAT has broad applications in other security areas, and we will further extend it along three directions. First, we will fully implement it and evaluate its efficacy on defending against diverse real-world privacy attacks launched by cloud providers. Second, we will augment the translator to automatically translate the big-data queries with access patterns on particular data into those without (e.g., oblivious executions). Third, we will further enhance DAE to support well isolated, secured operating system calls (e.g., library operating system calls [66]), so that MAAT will not only benefit big-data queries, but other distributed computing paradigms (e.g., graph queries [40]).

2.4 Research plan

This project will require two PhD students S1 and S2 to work for three years. In the first year, S1 will design and fully implement the KAKUTE system (part of **Objective 1**), and S2 will evaluate its performance and security strength on various real-world big-data queries (part of **Objective 1**). In the second year, S1 will use KAKUTE to fully develop the proposed Fine-grained Differential Privacy technique (part of **Objective 2**), and S2 will implement the algorithm proposed for this technique (part of **Objective 2**). In the third year, S1 will build the secure big-data compiler (part of **Objective 3**), and S2 will extensively study the privacy guarantee and performance of this compiler on real-world big-data frameworks (part of **Objective 3**).

Appendix (C): Non-text figures




Figure 1: The Reference Propagation technique with the given code (for **Objective 1**).



Figure 2: The Tag Sharing technique between fields in each record (for **Objective 1**).



Figure 3: The Data-locality-aware Asynchronous Enclave (DAE) call abstraction (for **Objective 3**).

References

- [1] GitHub - Azure/coco-framework. <https://github.com/Azure/coco-framework>.
- [2] K. A. V. P. M. Shaaban, and W. C.L. Heterogeneous computing: Challenges and opportunities. In *IEEE Computer*, 1993.
- [3] S. Akoush, L. Carata, R. Sohan, and A. Hopper. Mrlazy: Lazy runtime label propagation for mapreduce. In *Proceedings of the 6th USENIX Conference on Hot Topics in Cloud Computing*, HotCloud'14, pages 17–17, Berkeley, CA, USA, 2014. USENIX Association. URL <http://dl.acm.org/citation.cfm?id=2696535.2696552>.
- [4] C. S. Alliance. Top threats to cloud computing v1.0. <https://cloudsecurityalliance.org/topthreats/csathreats.v1.0.pdf>, Feb. 2010.
- [5] C. S. Alliance. The notorious nine: Cloud computing top threats in 2013. https://downloads.cloudsecurityalliance.org/initiatives/top_threats/The_Notorious_Nine_Cloud_Computing_Top_Threats_in_2013.pdf, Feb. 2013.
- [6] M. Armbrust, R. S. Xin, C. Lian, Y. Huai, D. Liu, J. K. Bradley, X. Meng, T. Kaftan, M. J. Franklin, A. Ghodsi, et al. Spark sql: Relational data processing in spark. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*, pages 1383–1394. ACM, 2015.
- [7] S. Bajaj and R. Sion. Trusteddb: A trusted hardware based database with privacy and data confidentiality. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data*, SIGMOD '11, pages 205–216, New York, NY, USA, 2011. ACM. ISBN 978-1-4503-0661-4. doi: 10.1145/1989323.1989346. URL <http://doi.acm.org/10.1145/1989323.1989346>.
- [8] J. Bell and G. Kaiser. Phosphor: Illuminating dynamic data flow in commodity jvms. In *Proceedings of the 2014 ACM International Conference on Object Oriented Programming Systems Languages & Applications*, OOPSLA '14, pages 83–101, New York, NY, USA, 2014. ACM. ISBN 978-1-4503-2585-1. doi: 10.1145/2660193.2660212. URL <http://doi.acm.org/10.1145/2660193.2660212>.
- [9] C. B.W.L., W. C.L., and K. Hwang. A migrating-home protocol for implementing scope consistency model on a cluster of workstations. In *PDPTA*, 1999.
- [10] Z. Chothia, J. Liagouris, F. McSherry, and T. Roscoe. Explaining outputs in modern data analytics. *Proceedings of the VLDB Endowment*, 9(12):1137–1148, 2016.
- [11] H. Cui, J. Wu, C.-C. Tsai, and J. Yang. Stable deterministic multithreading through schedule memoization. In *Proceedings of the Ninth Symposium on Operating Systems Design and Implementation (OSDI '10)*, Oct. 2010.
- [12] H. Cui, J. Wu, J. Gallagher, H. Guo, and J. Yang. Efficient deterministic multithreading through schedule relaxation. In *Proceedings of the 23rd ACM Symposium on Operating Systems Principles (SOSP '11)*, pages 337–351, Oct. 2011.
- [13] H. Cui, J. Simsa, Y.-H. Lin, H. Li, B. Blum, X. Xu, J. Yang, G. A. Gibson, and R. E. Bryant. Parrot: a practical runtime for deterministic, stable, and reliable threads. In *Proceedings of the 24th ACM Symposium on Operating Systems Principles (SOSP '13)*, Nov. 2013.
- [14] H. Cui, R. Gu, C. Liu, and J. Yang. Paxos made transparent. In *Proceedings of the 25th ACM Symposium on Operating Systems Principles (SOSP '15)*, Oct. 2015.
- [15] A. Dave and M. Zaharia. Arthur: Rich post-facto debugging for production analytics applications.
- [16] J. Dean and S. Ghemawat. Mapreduce: simplified data processing on large clusters. In *OSDI'04: Proceedings of the 6th conference on Symposium on Operating Systems Design & Implementation*, pages 10–10, 2004.
- [17] C. Dwork. Differential privacy. *Lecture Notes in Computer Science*, 26(2):1–12, 2006.
- [18] C. Dwork. A firm foundation for private data analysis. *Communications of The ACM*, 54(1):86–95, 2011.
- [19] C. Dwork, F. McSherry, K. Nissim, and A. Smith. Calibrating noise to sensitivity in private data analysis. In *TCC*, volume 3876, pages 265–284. Springer, 2006.
- [20] C. Dwork, F. McSherry, K. Nissim, and A. Smith. Calibrating noise to sensitivity in private data analysis. In *Proceedings of the Third Conference on Theory of Cryptography*, TCC'06, pages 265–284, Berlin, Heidelberg, 2006. Springer-Verlag. ISBN 3-540-32731-2, 978-3-540-32731-8. doi: 10.1007/11681878_14. URL http://dx.doi.org/10.1007/11681878_14.
- [21] W. Enck, P. Gilbert, B.-G. Chun, L. P. Cox, J. Jung, P. McDaniel, and A. N. Sheth. TaintDroid: an information-flow tracking system for realtime privacy monitoring on smartphones. In *Proceedings of the Ninth Symposium on Operating Systems Design and Implementation (OSDI '10)*, pages 1–6, 2010.
- [22] L. Feng, L. F.C.M., C. Heming, and W. Cho-Li. Confluence: Speeding up iterative distributed operations by key-dependency-aware partitioning. In *IEEE Transactions on Parallel and Distributed Systems (TPDS)*, 2017.
- [23] T. Frassetto, D. Gens, C. Liebchen, and A.-R. Sadeghi. Jitguard: Hardening just-in-time compilers with sgx. 2017.
- [24] C. Gentry, S. Halevi, and N. P. Smart. Homomorphic evaluation of the aes circuit. In *Advances in Cryptology-CRYPTO 2012*, pages 850–867. Springer, 2012.
- [25] C. Gentry et al. Fully homomorphic encryption using ideal lattices. In *STOC*, volume 9, pages 169–178, 2009.
- [26] M. A. Gulzar, M. Interlandi, S. Yoo, S. D. Tetali, T. Condie, T. Millstein, and M. Kim. Bigdebug: Debugging primitives for interactive big data processing in spark. In *Proceedings of the 38th International Conference on Software Engineering*, ICSE '16, pages 784–795, New York, NY, USA, 2016. ACM. ISBN 978-1-4503-3900-1. doi: 10.1145/2884781.2884813. URL <http://doi.acm.org/10.1145/2884781.2884813>.

- [27] Hadoop. Hadoop. <http://hadoop.apache.org/core/>.
- [28] X. Hu, M. Yuan, J. Yao, Y. Deng, L. Chen, Q. Yang, H. Guan, and J. Zeng. Differential privacy in telco big data platform. *Proceedings of the VLDB Endowment*, 8(12):1692–1703, 2015.
- [29] Intel. Software guard extensions programming reference. <https://software.intel.com/sites/default/files/329298-001.pdf>.
- [30] M. Interlandi, K. Shah, S. D. Tetali, M. A. Gulzar, S. Yoo, M. Kim, T. Millstein, and T. Condie. Titian: Data provenance support in spark. *Proc. VLDB Endow.*, 9(3):216–227, Nov. 2015. ISSN 2150-8097. doi: 10.14778/2850583.2850595. URL <http://dx.doi.org/10.14778/2850583.2850595>.
- [31] J. Jianyu, Z. Shixiong, A. Danish, W. Yuexuan, C. Heming, L. Feng, and G. Zhaoquan. Kakute: A precise, unified information flow analysis system for big-data security. In *Proceedings of the Annual Computer Security Applications Conference (ACSAC '17)*, 2017.
- [32] S. D. I. John and A. O. Osonde. Privacy preservation in the age of big data. In *RAND '16*, 2016.
- [33] Z. Jorgensen, T. Yu, and G. Cormode. Conservative or liberal? personalized differential privacy. In *Data Engineering (ICDE), 2015 IEEE 31st International Conference on*, pages 1023–1034. IEEE, 2015.
- [34] H. K., J. H., C. E., W. C.L., and X. Z. Designing ssi clusters with hierarchical checkpointing and single i/o space. In *IEEE Concurrency*, 1999.
- [35] M. Kazim and S. Y. Zhu. A survey on top security threats in cloud computing. *International Journal of Advanced Computer Science and Applications (IJACSA)*, 2015.
- [36] V. P. Kemerlis, G. Portokalidis, K. Jee, and A. D. Keromytis. Libdft: Practical dynamic data flow tracking for commodity systems. In *Proceedings of the 8th ACM SIGPLAN/SIGOPS Conference on Virtual Execution Environments, VEE '12*, pages 121–132, New York, NY, USA, 2012. ACM. ISBN 978-1-4503-1176-2. doi: 10.1145/2151024.2151042. URL <http://doi.acm.org/10.1145/2151024.2151042>.
- [37] L. King-Tin, S. Jinghao, H. Dominic, W. Cho-Li, L. Zhiquan, Z. Wangbin, and Y. Youliang. Rhymes: A shared virtual memory system for non-coherent tiled many-core architectures. In *ICPADS 2014*, 2014.
- [38] D. Logothetis, S. De, and K. Yocum. Scalable lineage capture for debugging disc analytics. In *Proceedings of the 4th annual Symposium on Cloud Computing*, page 17. ACM, 2013.
- [39] A. Machanavajjhala, J. Gehrke, D. Kifer, and M. Venkitasubramaniam. l-diversity: Privacy beyond k-anonymity. In *Data Engineering, 2006. ICDE'06. Proceedings of the 22nd International Conference on*, pages 24–24. IEEE, 2006.
- [40] G. Malewicz, M. H. Austern, A. J. Bik, J. C. Dehnert, I. Horn, N. Leiser, and G. Czajkowski. Pregel: a system for large-scale graph processing. In *Proceedings of the 2010 ACM SIGMOD International Conference on Management of data*, pages 135–146. ACM, 2010.
- [41] F. McSherry. Privacy integrated queries. In *Proceedings of the 2009 ACM SIGMOD International Conference on Management of Data (SIGMOD)*. Association for Computing Machinery, Inc., June 2009. URL <https://www.microsoft.com/en-us/research/publication/privacy-integrated-queries/>.
- [42] M. M.J.M., W. C.L., and L. F.C.M. Jessica: Java-enabled single-system-image computing architecture. In *Journal of Parallel and Distributed Computing*, 2000.
- [43] P. Mohan, A. Thakurta, E. Shi, D. Song, and D. Culler. Gupt: Privacy preserving data analysis made easy. In *Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data, SIGMOD '12*, pages 349–360, New York, NY, USA, 2012. ACM. ISBN 978-1-4503-1247-9. doi: 10.1145/2213836.2213876. URL <http://doi.acm.org/10.1145/2213836.2213876>.
- [44] L. Nelson, H. Sigurbjarnarson, K. Zhang, D. Johnson, J. Bornholt, E. Torlak, and X. Wang. Hyperkernel: Push-button verification of an os kernel. In *Proceedings of the 26th Symposium on Operating Systems Principles (SOSP '17)*.
- [45] J. Newsome and D. Song. Dynamic taint analysis for automatic detection, analysis, and signature generation of exploits on commodity software. 2005.
- [46] O. Ohrimenko, F. Schuster, C. Fournet, A. Mehta, S. Nowozin, K. Vaswani, and M. Costa. Oblivious multi-party machine learning on trusted processors. In *25th USENIX Security Symposium (USENIX Security 16)*, pages 619–636, Austin, TX, 2016. USENIX Association. ISBN 978-1-931971-32-4. URL <https://www.usenix.org/conference/usenixsecurity16/technical-sessions/presentation/ohrimenko>.
- [47] C. Olston, B. Reed, U. Srivastava, R. Kumar, and A. Tomkins. Pig latin: a not-so-foreign language for data processing. In *Proceedings of the 2008 ACM SIGMOD international conference on Management of data*, pages 1099–1110. ACM, 2008.
- [48] P. Paillier et al. Public-key cryptosystems based on composite degree residuosity classes. In *Eurocrypt*, volume 99, pages 223–238. Springer, 1999.
- [49] A. Papadimitriou, R. Bhagwan, N. Chandran, R. Ramjee, A. Haeberlen, H. Singh, A. Modi, and S. Badrinarayanan. Big data analytics over encrypted datasets with seabed. In *OSDI*, pages 587–602, 2016.
- [50] V. Pappas, V. P. Kemerlis, A. Zavou, M. Polychronakis, and A. D. Keromytis. Cloudfence: Data flow tracking as a cloud service. In *Proceedings of the 16th International Symposium on Research in Attacks, Intrusions, and Defenses - Volume 8145, RAID 2013*, pages 411–431, New York, NY, USA, 2013. Springer-Verlag New York, Inc. ISBN 978-3-642-41283-7. doi: 10.1007/978-3-642-41284-4_21. URL http://dx.doi.org/10.1007/978-3-642-41284-4_21.
- [51] pigmix. <https://cwiki.apache.org/confluence/display/PIG/PigMix>.

- [52] R. A. Popa, C. Redfield, N. Zeldovich, and H. Balakrishnan. Cryptodb: protecting confidentiality with encrypted query processing. In *Proceedings of the Twenty-Third ACM Symposium on Operating Systems Principles*, pages 85–100. ACM, 2011.
- [53] I. Roy, S. T. V. Setty, A. Kilzer, V. Shmatikov, and E. Witchel. Airavat: Security and privacy for mapreduce. In *Proceedings of the 7th USENIX Conference on Networked Systems Design and Implementation*, NSDI’10, pages 20–20, Berkeley, CA, USA, 2010. USENIX Association. URL <http://dl.acm.org/citation.cfm?id=1855711.1855731>.
- [54] F. Schuster, M. Costa, C. Fournet, C. Gkantsidis, M. Peinado, G. Mainar-Ruiz, and M. Russinovich. Vc3: Trustworthy data analytics in the cloud using sgx. In *Security and Privacy (SP), 2015 IEEE Symposium on*, pages 38–54. IEEE, 2015.
- [55] F. Shaon, M. Kantarcioglu, Z. Lin, and L. Khan. Sgx-bigmatrix: A practical encrypted data analytic framework with trusted processors. In *Proceedings of the 17th ACM conference on Computer and communications security (CCS ’10)*, 2017.
- [56] D. Sheng and W. Cho-Li. Error-tolerant resource allocation and payment minimization for cloud system. In *IEEE Transactions on Parallel and Distributed Systems (TPDS)*, 2013.
- [57] D. Sheng, R. Yves, V. Frederic, K. Derrick, W. Cho-Li, and C. Franck. Optimization of cloud task processing with checkpoint-restart mechanism. In *SC ’13*, 2013.
- [58] D. Sheng, K. Derrick, and W. Cho-Li. Optimization of composite cloud service processing with virtual machines. In *IEEE Transactions on Computers*, 2014.
- [59] A. Smith. Efficient, differentially private point estimators. *arXiv preprint arXiv:0809.4794*, 2008.
- [60] A. Smith. Privacy-preserving statistical estimation with optimal convergence rates. In *Proceedings of the Forty-third Annual ACM Symposium on Theory of Computing*, STOC ’11, pages 813–822, New York, NY, USA, 2011. ACM. ISBN 978-1-4503-0691-1. doi: 10.1145/1993636.1993743. URL <http://doi.acm.org/10.1145/1993636.1993743>.
- [61] Spark example. <https://spark.apache.org/examples.html>.
- [62] J. J. Stephen, S. Savvides, R. Seidel, and P. Eugster. Practical confidentiality preserving big data analysis. In *6th USENIX Workshop on Hot Topics in Cloud Computing (HotCloud 14)*, Philadelphia, PA, 2014. USENIX Association. URL <https://www.usenix.org/conference/hotcloud14/workshop-program/presentation/stephen>.
- [63] L. Sweeney. k-anonymity: A model for protecting privacy. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*, 10(05):557–570, 2002.
- [64] S. Technology. 7 most infamous cloud security breaches. <https://www.storagecraft.com/blog/7-infamous-cloud-security-breaches/>, Feb. 2017.
- [65] S. D. Tetali, M. Lesani, R. Majumdar, and T. Millstein. Mrcrypt: Static analysis for secure cloud computations. In *Proceedings of the 2013 ACM SIGPLAN International Conference on Object Oriented Programming Systems Languages & Applications, OOPSLA ’13*, pages 271–286, New York, NY, USA, 2013. ACM. ISBN 978-1-4503-2374-1. doi: 10.1145/2509136.2509554. URL <http://doi.acm.org/10.1145/2509136.2509554>.
- [66] C.-C. Tsai, D. E. Porter, and M. Vij. Graphene-sgx: A practical library os for unmodified applications on sgx. In *2017 USENIX Annual Technical Conference (USENIX ATC)*, 2017.
- [67] S. Tu, M. F. Kaashoek, S. Madden, and N. Zeldovich. Processing analytical queries over encrypted data. In *Proceedings of the VLDB Endowment*, volume 6, pages 289–300. VLDB Endowment, 2013.
- [68] Z. W. W. Cho-Li, , and L. F.C.M. Jessica2: A distributed java virtual machine with transparent thread migration support. In *IEEE Fourth International Conference on Cluster Computing (Cluster2002)*, 2002.
- [69] C. Wang, J. Yang, N. Yi, and H. Cui. Tripod: An efficient, highly-available cluster management system. In *Proceedings of the 7th ACM SIGOPS Asia-Pacific Workshop on Systems*, APSys ’16, 2016.
- [70] C. Wang, J. Jiang, X. Chen, N. Yi, and H. Cui. Apus: Fast and scalable paxos on rdma. In *Proceedings of the Eighteenth ACM Symposium on Cloud Computing*, pages 17–28. ACM, 2017.
- [71] N. World. Leaked icloud credentials obtained from third parties, apple says. <https://www.networkworld.com/article/3184471/security/leaked-icloud-credentials-obtained-from-third-parties-apple-says.html>, Feb. 2017.
- [72] J. Yang, H. Cui, J. Wu, Y. Tang, and G. Hu. Determinism is not enough: Making parallel programs reliable with stable multithreading. *Communications of the ACM*, 2014.
- [73] Y. Yu, M. Isard, D. Fetterly, M. Budiu, Ú. Erlingsson, P. K. Gunda, and J. Currey. Dryadlinq: A system for general-purpose distributed data-parallel computing using a high-level language.
- [74] M. Zaharia, M. Chowdhury, T. Das, A. Dave, J. Ma, M. McCauley, M. J. Franklin, S. Shenker, and I. Stoica. Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing. In *Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation*, pages 2–2. USENIX Association, 2012.
- [75] K. Zhang, X. Zhou, Y. Chen, X. Wang, and Y. Ruan. Sedic: privacy-aware data intensive computing on hybrid clouds. In *Proceedings of the 18th ACM conference on Computer and communications security*, pages 515–526. ACM, 2011.
- [76] W. Zheng, A. Dave, J. G. Beekman, R. A. Popa, J. E. Gonzalez, and I. Stoica. Opaque: An oblivious and encrypted distributed analytics platform. In *NSDI*, pages 283–298, 2017.
- [77] L. Zhiquan, L. King-Tin, W. Cho-Li, , and S. Jinshu. Powerock: Power modeling and flexible dynamic power management for many-core architectures. In *IEEE Systems Journal*, 2016.