

1 Research Background

This section introduces the background of big-data frameworks (§1.1), software-based privacy techniques (§1.2), hardware-based privacy techniques (§1.3), motivation of objectives (§1.4), and related work (§1.5 and §1.6).

1.1 Big-data Computing Frameworks

Big-data frameworks (e.g. Spark [64], DryadLINQ [63], MapReduce [14]) are popular for computations on tremendous amounts of data. These frameworks provide self-defined functions (e.g., map and reduce) to let computation providers write their algorithms to perform queries or transformations on data, and these frameworks automatically apply the functions on the data stored across computers in parallel. To exchange intermediate results across hosts, shuffles operations are often invoked and sometimes they are performance bottlenecks in big-data frameworks. For instance, Spark [64] has intensive data shuffling between its map and reduce stages.

To avoid excessive computation, many DISC systems adopt the lazy transformation approach [40, 63, 64]. Spark uses lazy transformations (e.g. MAP) for efficiency, and calls to these transformations only create a new data structure called RDD with LINEAGE (the sequence of transformation operations for a data record). The real transformations are only triggered when collecting operations (e.g. COLLECT, COUNT) are called. These collecting operations trigger transformations along lineages, where unnecessary computations are avoided. **Objective 1** (§2.1.2) leverages the lazy transformation feature in big-data frameworks to develop the new Reference Propagation technique.

1.2 Software-based Privacy Techniques

Data Flow Tracking (DFT) is a mandatory access control technique for preventing sensitive information leakage [38]. DFT attaches a tag to a variable (or object), and this tag will propagate throughout the computation at runtime. For example, a variable-level dataflow tracking will involve combinations of tags of two variables in each instruction, using an IOR operation. Different granularities of computation may incur different levels of computation overhead. Lower level (e.g. byte-level) tracking will consume a lot of resources, as each byte of data in an DFT system has its own tags [29]. DFT has been applied to various areas, such as preventing sensitive information (e.g. GPS data and contacts) leakage in cellphone [17, 56], providing secure web services [43] and server programs [29]. To the best of our knowledge, no DFT system exists for big-data computing.

Complimentary to DFT, statistical techniques allow the aggregation of sensitive data while preserving the privacy of any individual one. These statistical techniques include kanonymization methods [32, 55] or differential privacy [34, 37, 46], which enforce statistical bounds to prevent individual information leakage.

However, these statistical techniques are either not secure (kanonymization) or suffering from great losses of accuracy (differential privacy). A recent work [23] reported more than 30% losses of accuracy when the security guarantee is high (the probability of leakages is low). In such case, a simple KMeans program will return centroids far from the accurate ones, and the accuracy loss rate is much larger than the training error rate which is several percents in practice. Overall, despite much effort, existing differential privacy techniques can only favor privacy or utility of results, but not both, and key reason is that these techniques lack a precise tracking of how sensitive data fields flow to query results, so they have to take a coarse-grained approach, which conservatively adds noise to all fields and records.

In practice, only some fields in a data record may be sensitive and it is too rigorous to make all fields inaccurate. Moreover, different data records may have various security levels. For example, in an Taobao order record $\langle t, \text{userId}, \text{productId} \rangle$, the `userId` field is sensitive and it must not be leaked, and different owners of movie rating records may demand different levels of protection for their data. **Objective 2** (§2.2) proposes a novel fine-grained differential privacy technique, which combines the strengths of DFT and differential privacy.

1.3 Hardware-based Privacy Techniques

Trusted Execution Environment (TEE) is a promising technique for protecting computation in a public cloud even if the cloud’s operating systems or hypervisors are compromised. When a program runs in a TEE, the memory can not be seen by outside parties. Intel-SGX [25] is a popular commercial TEE implementation which runs programs in a hardware-protected enclave, so code and data are protected and can not be seen by the attackers. Compared with the approach of computing on encrypted data (§1.5), TEE is much safer and 100X to 1000X faster. A SGX-based work Opaque [66] reports only 30% overhead compared to native big-data queries, promising for real-world deployment.

However, to practically run Java big-data queries with SGX, two open challenges remain. First, running programs in an enclave needs rewriting the readily pervasive Java queries into SGX-compatible C++, which is extremely time-consuming and error-prone. Recent work [6, 66] run Zookeeper and SparkSQL in enclaves, and both of them rewrote codes running in enclaves using C++. In the long run, SGX tends not to support Java (partially due to minimizing TCB, see below), thus running unmodified big-data queries with SGX is highly desirable.

Second, existing SGX-based systems for big-data have too large Trusted Computing Base (TCB). Running a whole JVM in enclaves is considered impractical as size of TCB is too large [6], because JVM code comes from many different parties/vendors and extremely hard to verify. SGX-BigMatrix [48] proposes a secure and oblivious vectorization abstraction for Python, but it needs to rewrite programs with its abstraction. Overall, it is challenging to run programs in enclaves with minimum TCB. **Objective 3** (§2.3) aims to tackle these two open challenges.

1.4 Motivation of objectives

Data leakage is a top threat in cloud computing, and this threat becomes even more pronounced due to the pervasive big-data computing trend and more and more data providers, including application vendors individual computer users, are putting their important data on clouds. In a data provider’s perspective, both computational providers and cloud providers can breach the privacy of the data. Therefore, this proposal aims to preserve the data provider’s privacy by going two directions. First, we propose two novel complementary techniques in Objective 1 (Kakute) and Object 2 (fine-grained differential privacy) to avoid leakage from the computational providers. This two objectives can already effectively preserve privacy in a private cloud setting. Second, we propose a new privacy-preserving just-in-time compiler to avoid privacy leakage from the (public) cloud providers. By integrating the outcomes from all three objectives, data privacy will be effectively preserved against both the computational providers and cloud providers.

1.5 Related work by others

Computing on encrypted data. Homomorphic encryption [16, 21, 41] is a technique for performing computations on encrypted data in untrusted environments. Homomorphic encryption contains two kinds: Fully homomorphic encryption (FHE) and partial homomorphic encryption. Partial homomorphic encryption (e.g. Additive Homomorphic Encryption [41]) incurs a much lower overhead compared with FHE. A evaluation [20] on FHE shows a 10^9 slowdown, which is acceptable in practice. Systems that adopt PHE (e.g. Monomi [59], Crypsis [54], CryptDB [45], MrCrypt [57]) reports a much better overhead, but it has limited expressiveness (e.g., SQL operators) and requires extra trusted servers for computations. Seabed [42] proposes asymmetric encryption schemes and reduces performance overhead incurred by AHE, but it still has limited expressiveness.

Trusted execution environments (TEE). Intel SGX is a promising technique to provide privacy-preserving analytics in public clouds. Compared with software-based solutions, hardware-based solutions incur much lower overhead. TrustedDB [4] is a hardware-based secure database. VC3 [47] proposes a secure distributed analytic platform with read-write validations on Mapreduce [14]. Opaque [66] supports secure and oblivious SQL operators on SparkSQL [3]. However, all these systems have limited expressiveness (e.g. SQL

operators), and VC3 even needs to rewrite the program with C++. A recent work [39] proposes a oblivious machine learning framework on trusted processors. BigMatrix [48] proposes an oblivious and secure vectorization abstraction on python, but it has limited expressiveness and it needs to rewrite the original program with this new abstraction. Although BigMatrix provides guideline for writing a oblivious program, but it would be a time-consuming and error-prone process.

Big-data privacy systems. Big-data privacy has been a top emerging threat as more and more user sensitive data is captured and processed in clouds. Airavat[46], PINQ [34] and GUPT [37] propose to apply differential privacy [15, 35] in Mapreduce, to prevent leakage from user query, but differential privacy can result in incorrect results. Sedici [65] proposes to offload sensitive computations to private clouds. MrLazy [2] proposes a framework of combining data provenance and static DFT analysis for self-defined queries, to provide fine-grained information flow for security. However, static DFT is not precise and may suffer from false positive. KAKUTE provides fine-grained information control of sensitive data, with no need to modify the original program.

Data provenance. Given an output record, the data provenance technique can identify inputs that produce this output, so it has shown wide range of applications, including security and debugging [22, 68]. RAMP [24], Newt [31], Pig [19] and Titian [26] adopt a record-level tracking approach for data provenance in DISC frameworks. Chothia [8] introduces a novel framework for output explaining in iterative programs with differential dataflow abstraction. These systems also adopt the record-level tracking approach, so they have low precision in programs containing many-to-many transformations. KAKUTE, however, provides fine-grained DFT in terms of data field level, so KAKUTE is much more precise than existing these existing systems on data provenance.

1.6 Related work by the PI and co-I

The PI is an expert on secure and reliable distributed systems [9–12, 18, 27, 61, 62]. The PI’s works are published in top conferences on systems (OSDI, SOSP, SOCC, TPDS, and ACSAC) and programming languages (PLDI and ASPLOS). The co-I is an expert on high-performance computing [1, 7, 28, 36, 67], fault-tolerance [49, 50], and Java compilers [30, 51, 60]. The co-I’s works are published in top systems conferences (Cluster ’02, SC ’13, and ICPADS ’14) and journals (JPDC ’00, TPDS ’13, IEEE Tran. Computers ’14). As preliminary results for this proposal, the PI has published Kakute [27] in ACSAC ’17, and the PI and co-I have collaborated to publish Confluence [18] in TPDS ’17.

2 Research Plan and Methodology

This section first proposes the three objectives in this proposal.

2.1 Objective 1: preventing big-data leakage with Precise IFT

This section presents PAXOS performance problem (§2.1.1) and KAKUTE (§2.1.2).

2.1.1 Problem: existing IFT systems are slow and incomplete for big-data

No IFT system exists for big-data, and we attribute it to two major challenges. First, existing IFT systems incur high performance overhead, especially for data-intensive computations. We ran a recent IFT system Phosphor [5] in Spark with a WordCount algorithm on a dataset that is merely 200MB, and observed 128X longer computation time compared with the native Spark execution (§1.1). The second challenge is on the architecture of DISC frameworks. DISC frameworks usually contain shuffle procedures which redistribute data across hosts (DISC frameworks’ worker nodes). However, most existing IFT systems ignore dataflows across hosts. For the few [43] who support cross-host dataflows, transferring all tags in shuffles consumes excessive network bandwidth. Therefore, efficient cross-host tag propagation is crucial but missing in DISC.




Figure 1: Threat model).

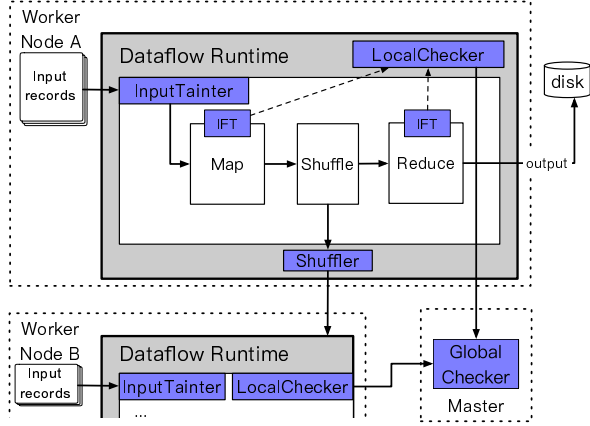


Figure 2: KAKUTE architecture.

2.1.2 KAKUTE: a fast, precise IFT system for big-data

This paper presents KAKUTE¹, the first precise and fine-grained information flow analysis system in DISC frameworks. Our key insight to address the IFT efficiency challenge is that multiple fields of a record often have the same tags. Leveraging this insight, we present two new techniques, Reference Propagation and Tag Sharing. Reference Propagation avoids unnecessary tag combinations by only keeping the *lineage of tags* in the same UDF, while Tag Sharing reduces memory usage by sharing tags among multiple fields in each record. To tackle the architecture challenge, KAKUTE completely captures dataflows in shuffles, and it efficiently reduces the amount of transferred IFT tags using Tag Sharing.

We will implement KAKUTE in Spark. We leverage Phosphor [5], an efficient IFT system working in the Java byte-code level. KAKUTE instruments computations of a Spark worker process to capture dataflow inside user-defined-functions (UDFs). Dataflow information is kept in a worker process and KAKUTE propagates it to other processes while shuffling. Therefore, IFT is completely captured across hosts and processes. In this paper, DISC frameworks and KAKUTE are trusted; UDFs are untrusted and they may be malicious or buggy. KAKUTE provides different granularities of tracking with two types of tags: INTEGER and OBJECT tags. INTEGER provides 32 distinct tags, which is suitable for detecting information leakage and performance bugs. OBJECT provides an arbitrary number of tags, which is suitable for data provenance and programming debugging. KAKUTE provides a unified API to tackle diverse problems. Based on this unified API, we implemented 4 built-in checkers for 4 security and reliability problems: sensitive information leakage, data provenance, programming and performance bugs.

Preliminary results. We evaluated KAKUTE on seven diverse algorithms, including three text processing algorithms WordCount [53], WordGrep [31] and TwitterHot [53], two graph algorithms TentativeClosure [53] and ConnectComponent [53], and two medical analysis programs MedicalSort [44] and MedicalGroup [44]. These algorithms cover all big-data algorithms evaluated in two related papers [22, 26]. We evaluated these algorithms with real-world datasets that are comparable with related systems [8, 22, 26]. We compared KAKUTE with Titian [26], a popular provenance system, on precision and performance. Our evaluation shows that: (1) KAKUTE is fast. Kakute had merely 32.3%




Figure 3: KAKUTE execution time normalized to Spark executions.

¹Kakute is a precise, multi-purpose weapon used by Ninja.

Algorithm 1: Naive Aggregation Model

Input: Dataset T, dataset size N, privacy budget ϵ_k for security level k, output range (min, max)
n = 4
for $i \leftarrow 1$ **to** n **do**
 $O_i \leftarrow f(T_i)$;
 if $O_i > \max$, $O_i \leftarrow \max$ if $O_i < \max$, $O_i \leftarrow \min$
for dimension j of the output O **do**
 $k \leftarrow \max(\text{getLevel}(O_j))$;
 $O_j \leftarrow \frac{1}{n} \sum_{i=1}^n O_{ij} + \text{Lap}(\frac{\max - \min}{\epsilon_k})$
Output: O

overhead with INTEGER tag, suitable for production runs; and (2) KAKUTE is precise; it effectively detected 13 real-world security and reliability bugs presented in other papers [13, 22, 46].

Future work. We will further improve the practicality of KAKUTE in two dimensions. First, we will extensively evaluate its performance on broad big-data programs. Second, we will leverage it to improve existing security techniques, including anonymization techniques and diff privacy (Objective 2).

2.2 Objective 2: fine-grained diff privacy

We are proposing to combine IFT and differential privacy. With IFT, security levels are propagating throughout computation, while differential privacy is for preventing individual information leakage. In this model, different records and fields can have various security levels, therefore, a general and fine-grained differential private framework can be achieved. Programmer does not need to modify their analysis programs and tradeoff between usability and security can be achieved with different users.

It needs to determine the relation of the privacy budget and the security level. In a simplest model, there are 5 security levels: insensitive, dp-1, dp-2, dp-3, non-released. Insensitive record can be release directly, while non-released can not be used in any computation to the final result. dp-1 to dp-3 varies in terms of their privacy budgets for differential privacy. We may also consider the (δ, ϵ) -differential privacy model.

To estimate the range of a output, so that the laps noise distribution can be determine by running the same, differentially private percentile estimation algorithm on the inputs to compute the 25-th and the 75-th percentile privately (a.k.a, lower and upper quartiles) for the inputs.

We extend the differential privacy model developed in previous work [52]. We introduce two noise aggregation model that make use the fact that different record may have different security level (protection level).

A simplest model is to use the ϵ inferred by the highest security level, then add noise to the output directly. Add noise to the output with $f(x) + \text{Lap}(\frac{\max - \min}{\epsilon_{\max}})$. This simple model may generate too much noise to the final result, which make the result unusable. The complete algorithm is showed in Algorithm 1.

In this model, data is partitioned into multiple parts (size of each part is m). Each partition may consists different level. The noise aggregator adds noise to final result of each partition according to the security level of each partition.

Formally, data is partitioned into multiple parts denoted as p_1, p_2, \dots, p_n . The security level and its corresponding ϵ are $\epsilon_1, \epsilon_2, \dots, \epsilon_n$. Therefore, the final aggregated result is

Algorithm 2: Combined Aggregation Model

Input: Dataset T, dataset size N, privacy budget ϵ_k for security level k, output range (min, max)

$n = 4$

for $i \leftarrow 1$ **to** n **do**

$O_i \leftarrow f(T_i)$;

if $O_i > \max$, $O_i \leftarrow \max$ **if** $O_i < \max$, $O_i \leftarrow \min$ **for** *dimension j of the output O* **do**

$k \leftarrow \max(\text{getLevel}(O_{ij}))$;

$O_{ij} \leftarrow O_{ij} + \text{Lap}(\frac{\max - \min}{\epsilon_k})$

foreach *partition i of the output O* **do** $O_j \leftarrow \frac{1}{n} \sum_{i=1}^n O_{ij}$;

Output: O

$$\frac{\sum_{i=1}^n f(x_i) + \text{Lap}(\frac{\max_i - \min_i}{\epsilon_i})}{n} \quad (1)$$

The complete algorithm is showed in Algorithm 2.

The error of the final result incurs in this aggregator comes from two parts: the Laps noise error and the partition error. It is crucial to reduce the final error incurs by this aggregator while keeping the differential security guarantee. The error of this model (when the eps is the same) is

$$|\frac{1}{n} \sum_{i=1}^n f(T_i) - f(T)| + \frac{1}{n} \sqrt{2} \Delta f \sum_{i=1}^n \frac{1}{\epsilon_i} \quad (2)$$

, and it should be minimized.

As final result of each partition may consist different security levels, the final noise can be less than previous method.

If we can get the statistic of security levels, we can estimate the block size in a better way. In specific, suppose there are k security levels and M_k means that M_k partitions has a maximum security level as k. The the above equation can be rewritten as

$$|\frac{1}{n} \sum_{i=1}^n f(T_i) - f(T)| + \frac{1}{n} \sqrt{2} \Delta f \sum_{i=1}^k \frac{1}{\epsilon_i} * M_i \quad (3)$$

M_k can be calculated according to the total partition number n and the number of block that has security level as k (denoted as p_k), then we can further formulate it as:

2.3 Objective 3: building a secure just-in-time compiler

In the project, we plan to tackle the problem above. We are proposing to run unmodified Java program in enclaves to protect computation in public clouds or untrusted servers. We will design a Just-In-Time compiler for JVM and run secured functions in enclaves.

Thread Model We evaluate the program setup in public clouds. In a public cloud, only data provider, and a portion of code in the analytic platform that is running in the enclaved are trusted. In specific, the JIT compiler and the secure functions should be trusted. All other components, including operating system, hypervisor are trusted. Figure 4 shows the model (TODO).




Figure 4: Threat model in public cloud. All grey boxes are not trusted and red lines represent potential leakages or attacks. Data is encrypted outside enclave and is in blue.

Figure 5 shows the architecture of the system. Programmers annotate some functions as secure, so code of the functions and data processed by these functions should be kept as secrets. The annotated secure functions are compiled and executed inside enclaves so that data and code will not be leaked. The Java byte-codes of these function are compiled to native enclave codes, and are executed inside enclaves upon function calls.

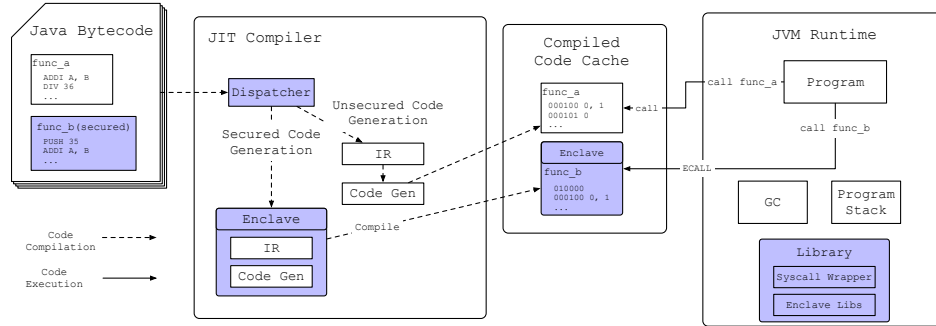


Figure 5: Architecture of Running Java with SGX. Some of the functions are annotated as secured function and they are compiled by a JIT in a enclave. The compiler compiles the functions into enclave code, so that the program will call into a enclave when calling into this function.

There are two challenges that we need to address in our design: running unmodified Java programs with minimum TCB and reducing excessive enclave transitions.

A straightforward approach to run unmodified Java programs in enclaves for code and data protections is to run the whole JVM in enclaves. However, as argued in a previous work [6], it will blow up the TCB and cause a high overhead by running the whole JVM in enclaves.

SGX is for protecting data and code running in enclaves, but code inside enclaves also have accesses to the regular memory region. Therefore, code in the enclaves can write sensitive data to unprotected memory regions when the code is compromised. The OpenJDK implementation of JVM contains up to millions lines of code, which is impractical for formal verifications as it is time-consuming with current verification methods (§1.5).

Intel SGX contains a protected memory region call Enclave Page Cache (EPC), and evictions from this cache cause expensive encryption costs. Although next generation of SGX [33] may support a larger

EPC, current generation of SGX only has an EPC up to 128MB. In practice, only around 90 MB can be allocated. Therefore, running programs with large memory consumption will cause much higher overhead. A recent work [66] shows that running program below this limit incur an only 7.46% overhead, while slightly exceeding this limit causes 50% to 60% overhead. Therefore, we have to reduce the TCB size for running unmodified Java programs in enclaves.

To reduce TCB size, we propose to run only some programmer-annotated functions in enclave. To this end, we propose a split execution framework. Only the Intermediate Representation (IR) and Code Generation model are trusted in the system. There will be two instances of compilers for compiling Java bytecode, one for secure code compilations and one for untrusted code compilations. Therefore, the TCB is greatly reduced, as we do not need to trust the JVM runtime (e.g. Garbage Collection). Also, as only some functions should be executed securely, the split execution framework should not cause high performance overhead.

When calling into and out of a function in enclaves, an ECALL and OCALL will be invoked in the CPU. In specific, the CPU is trapped into a new mode (except from the User and Privileged mode), and the current frame is encrypted and saved. A recent work [58] shows that enclave transitions are 60X slower than system calls and several hundreds times slower than user function calls. Moreover, encryption and decryption are required for secure function calls, which makes enclave transitions even more time-consuming. Our evaluation on a recent privacy-preserving data analytic system [66] shows that it incurs 3.4k transitions for processing 10k data (for two operations select and groupBy). In fact, these processing functions can be pipelined and processed in a single enclave function. Our project takes reducing transitions as one of our targets.

We introduce two techniques, Cost-based Compilation and Asynchronous Enclave Call, to tackle this challenge. Cost-based Compilation can make use of the JVM hotspot features, and analyse the hotspot enclave functions. It builds a tree of callers and callee of functions, and combines two enclaves if the marginal benefit of combining them is larger than the transition cost. Initially, only function b and d are running in enclaves, but the compiler finds out that cost of running c in an enclave is less than the transition cost (assume it to be 2), then the whole function a will be compile as an enclave. In another case where running function c takes a high cost, combinations of enclaves will not happen. Therefore, transition of enclaves can be reduced. This technique can be applied online or offline. In a offline version, it sample the program and finish the optimisation offline.

Asynchronous Enclave Call convert the synchronous enclave calls to asynchronous enclave calls. In specific, when a secure function is called, the function call and its parameters are put into a QUEUE which will be fetch by the enclave can be executed. After finishing execution, the result will be stored in QUEUE_R and the current execution will resume and continue. This technique makes use of the multi-core hardware architecture and enclaves and the main program are running in separate threads.

These two technique can be applied simultaneously, the compiler can run the sample and optimisation offline. After that, the compiler run a enclave and the program in separate threads to avoid transition of enclaves.

2.4 Research timeline

This project will require two PhD students S1 and S2 to work for three years. In the first year, S1 will design and fully implement the KAKUTE system (part of **Objective 1**), and S2 will evaluate its performance and robustness on various real-world big-data queries (part of **Objective 1**). In the second year, S1 will use KAKUTE to fully develop the proposed fine-grained privacy model (part of **Objective 2**), and S2 will implement the two algorithms proposed for this model (part of **Objective 2**). In the third year, S1 will build the secure big-data compiler **Objective 3** and S2 will evaluate this compiler on diverse real-world big-data queries.